

# MACHINE LEARNING FOR TRADING

Diana Yau

## PROBLEM

- Given an in-sample period of a particular stock, how can we determine whether to buy, hold, or sell on an out-of-sample period using reinforcement-based learning?
  - In-sample period: January 1, 2008 to December 31, 2009.
  - Out-of-sample period: January 1, 2010 to December 31, 2011.

## SOLUTION


- Language: Python
  - NumPy and Pandas available as open-source.
  - Easy to access data.
- Q-Learning where  $Q[s,a]$  is a 2D table stores the consequent rewards where 's' is the state and 'a' is the action.
  - Using combined and normalized indicators as states, query the table to get the new Qvalue
- Return a trades DataFrame (a table where the dates listed in the out-of-sample period correspond to a state).

## IMPLEMENTATION – INDICATORS

- 3 indicators: Price/SMA, Bollinger Bands® %B, Momentum
  - Using indicator values of: Buy: 2, Hold: 1, Sell: 0.
  - Gives the maximum number of states as 222 and minimum as 000 with 3 actions.
  - For any day where there are 2 or more indicator values to buy/sell, consider there is a signal.
  - Consider having a window size of 10 days.

# IMPLEMENTATION - QLEARNING


- Q-learning based strategy using a Q-Learner.
  - Not greedy because it takes the action at the current time.
- Policy  $\pi(s) = \operatorname{argmax}(Q[s, a])$
- $Q[s, a]$  is a 2D table storing all known states and actions.
- A tuple of the state( $s$ ), current action( $a$ ), new state ( $s'$ ), and reward( $r$ ) is passed iteratively.




## IMPLEMENTATION – QLEARNING (CONT.)

$$Q'[s, a] = (1 - \alpha)Q[s, a] + \alpha(r + \lambda * Q[s', \operatorname{argmax}(Q[s', a'])])$$

  
New Q value

  
old value

  
discounted reward for all future actions

  
new best estimate

$\alpha$ : learning rate (0 to 1), typically 0.2

$\lambda$ : discount rate (0 to 1)

$r$ : immediate reward

## IMPLEMENTATION – TRAINING DATA

while not converged:

    x = indicators

    querySetState(x)

        for each day:

            r = calculate reward

            a = query(x,r)

            add action to dataframe of trades

            x = new state

check if converged

## CHALLENGES – FEATURE SELECTION

- Which indicators to use? What is a good number of indicators?
  - Price/SMA (current price/simple mean average)
  - Price/EMA (current price/exponential moving average)
  - Bollinger Bands® %B ( $[\text{current price} - \text{lower bound}] / [\text{upper bound} - \text{lower bound}]$ )
    - Bollinger Bands® ( $\text{SMA} \pm 2 \text{ Standard Deviations}$ )
  - Momentum (current price / previous price by a specific window of days ago)
  - %Price (difference between two moving averages)
- Lagging indicators
  - Backfilled data by the window size



## CHALLENGES (CONT.)

- How can the indicator be interpreted as buy or sell signals?
  - Suggestions: Buy: 1, Hold: 0, Sell: -1 or Buy: 2, Hold: 1, Sell: 0
- How can a combination of these indicators tell us what action to take?
  - How many indicators need to be validated to count the day as a buy/sell signal?
  - Considering the indicators as states?
- Trial and error: adjusting the learning rate and discount rate to maximize more trades
  - More trades = more potential to maximize holdings

## FUTURE WORK

- Using a set cash value, compare the cumulative returns, mean and standard deviation of daily returns
  - Finding the maximum reward output by adjusting the alpha (learning rate) and lambda (discount rate).
- Comparing with a benchmark (S&P 500) and a classification-based learner, which can be a Random Forest learner.
  - Things to consider: Setting leaf size to avoid overfitting in-sample period
- Increasing the number of indicators to solidify position (e.g., RSI – relative strength index, stochastic indicator, etc.)
  - Different trading setups: Stops, trailing stops, stop-loss
- Wider in-sample period with considerations of special days (e.g., stock split forward/backward, fed rates hike, etc.)

**Q&A**